

Discrete Choice Model Estimation with R

Dan Yavorsky

Table of contents

Welcome	4
Prerequisites	5
Acknowledgments	6
Colophon	6
I Introduction	7
1 Introduction	8
1.1 Recap of Train Ch. 1	8
1.2 A Simple Example	9
1.3 A Binary Logit Model Example	10
1.4 Key Learnings	12
2 Properties of Discrete Choice Models	13
II Behavioral Models	14
3 Logit	15
4 GEV	16
5 Probit	17
6 Mixed Logit	18
7 Variation on a Theme	19
8 Numerical Maximization	20
III Estimation	21
9 Drawing from Densities	22
10 Simulation-Assisted Estimation	23

11 Individual-Level Parameters	24
12 Bayesian Procedures	25
13 Endogeneity	26
14 EM Algorithms	27
References	28
 Appendices	 29
A Appendix	29

Welcome

Warning

This is an early draft. The book is incomplete. Every aspect of it is subject to change. Many things might be incorrect. Are you sure you want to read this yet?

You don't understand it, until you have coded it.

It's a mantra that my mentors instilled in me in graduate school, and one that I propagate on my students today. I believe it in deeply. If you want to understand a statistical model, it is insufficient to interact with the model only through pen and paper. No matter how well you organize the subscripts or how masterfully you interweave elements of the Greek and Roman alphabets, you won't really "get it." To deeply understand a statistical model, you need to be able to simulate data from the model, and then take that simulated data to an estimation routine that enables you to recover the parameters of interest.

This is something you learn to do in graduate school. But, unless you have one of the most caring and pedagogical advisers, no one teaches it to you! Instead, it's a skill that students develop independently and inefficiently between classes and assignment due dates or, in my case, after I was knee-deep in my dissertation research. There is a tremendous imbalance between how immensely important this skill is and the time and instruction dedicated to it.¹ That imbalance is the motivation for this book.

I address this gap in the specific domain of Discrete Choice Modeling. I do so alongside Kenneth Train's masterful text *Discrete Choice Methods with Simulation* (Second Edition) freely available online at <https://eml.berkeley.edu/books/choice2.html>.

Train and I align in our thinking:

"...the true value of [] choice modeling is the ability to create tailor-made models. The computation and programming steps that are needed to implement a new model are usually not difficult. An important goal [] is to teach these skills as an integral part of the exposition of the models themselves."²

¹In this context, a *closed-form expression* means a way of writing the integral so that the anti-derivative sign is not part of solution. For example, the integral $\int x dx$ has the closed for expression $x^2/2$ plus some constant. We will see later that the Extreme Value distribution is often chosen for f predominantly because it leads to a closed for expression for the choice probability $p(y|x)$.

²More precisely, $p(y = 2|x) = \Pr(x + \varepsilon > 1|x = 0.5) = \Pr(\varepsilon > 0.5) = \int_{0.5}^1 f(\varepsilon) d\varepsilon = (0.5\varepsilon)|_{0.5}^1 = 0.25$.

What is “not that difficult” for an experienced researcher can be immensely difficult for a typical graduate student. The aim of this book is to ease that difficulty. I demonstrate R code alongside the math. You will see the mathematical derivations and pseudo-code in Train (2009) transformed into working R code. No mysteries will remain. A goal, in fact, is to provide you with much of the same instruction that Train’s students received when taking his course. To borrow a quote from Quintilian, Cobbett, Cooke, and many others: my goal is to take you through the process of programming the simulation and estimation of discrete choice models *not so that you can understand, but so that you cannot possibly misunderstand*.

Prerequisites

This book is intended for a narrow audience, predominantly graduate students with an interest in discrete choice modelling who will find value from seeing and interacting with the programmatic implementation of the multinomial logit and its extended family of related models. In other words, someone who read Train (2009) and thought *how would I code that?*

We will simulate data from the statistical models and then estimate the parameters of those models from the simulated data. And in fact, estimation of many models will require approximating one or more integrals through a second form of simulation (the one referenced in the title of Train’s book). We will do all of this in R, a freely available software environment for statistical computing and graphics. As a result, I assume you are reasonably familiar with R. If not, there is a tremendous set of free online R resources collected and organized at <https://www.bigbookofr.com/>. Popular books include Wickham, Cetinka-Rundel, and Grolemund (2023) and Wickham (2019).

I also assume you have taken introductory statistics or econometrics courses, as the concepts and techniques taught there are foundational for understanding and estimating the discrete choice models covered by Train (2009). In particular, you should have no uncertainty about the difference between a model, an estimator, and an estimate. To briefly review:

- a *model* is the set of mathematical assumptions about how data are generated,
- an *estimator* (or equivalently, the estimation routine) is an algorithm or function of the data, and
- an *estimate* is the result of applying the estimator to a particular dataset.

In my experience, students are often given a model and an estimator, after which much time in the classroom is spent deriving properties of the estimator for that particular model. Often a homework assignment follows which asks students to implement the estimator on a dataset to find an estimate. This approach puts almost no emphasis on the specification of the model or the choice of estimator. For example, should part of the model be specified as $\beta_0 + \beta_1 x_1$ or should it be $\beta_0 + \beta_1 x_1 + \beta_2 x_2$? And once we have specified a model, should we derive an estimator via least squares, the method of moments, maximum likelihood, a Bayesian approach, or some other way?

If my description captures your experience in introductory statistics and econometrics courses and you would like to review key ideas, I highly recommend [Abramovich and Ritov \(2023\)](#) on the topic of mathematical statistics, [Goldberger \(1991\)](#) and [Kennedy \(2008\)](#) for econometrics, and [McElreath \(2018\)](#) and [Gelman et al. \(2013\)](#) for Bayesian statistics. In this book, we will use Maximum Likelihood and Bayesian methods to estimate the parameters of the models discussed in the forthcoming chapters.

Acknowledgments

I am immensely grateful to the teams that work on the R-Project, those at Posit who provide RStudio, Positron, and Quarto, and the related communities of developers, academics, and R users. The free tools they provide are exceptional and I use them daily. The welcoming communities they have established provide the necessary support humans need, be it technical troubleshooting or emotional encouragement.

I also thank my academic mentors Ella Honka, Peter Rossi, and Eric Bradlow, and (although we have yet to meet) Kenneth Train as well as the folks that tolerate me through my consulting work, including Prachi Bhalerao, June Wu, Chris Diener, Keith Chrzan, and the hosts and attendees of Sawtooth Software’s annual Analytics and Insights Summit. I have learned so much from these researchers’ guidance and often their written work.

Special thanks go to generous people who reviewed and assisted with drafts of this book, including Darren Aeillo, Kalyan Rallabandi, and Geoff Zheng.

Unquestionably, my deepest thanks go to Alison, Zach, and Ben for their patience and support.

Colophon

An online version of this book is available at <https://dcms-r.danyavorsky.com>. The source of the book is available at <https://github.com/dyavorsky/dcms-r>. The book is authored using [Quarto](#), an open-source scientific and technical publishing system that makes it easy to create articles, presentations, websites, books, and other publications that combine text and executable code.

Part I

Introduction

1 Introduction

It is instructive to start with the first sentence of Train (2009) Section 1.3, “Discrete choice analysis consists of two interrelated tasks: **specification** of the behavioral model and **estimation** of the parameters of that model” (emphasis added). In particular, chapters 1–7 of Train (2009) only specify models; there is not even a hint of model estimation. And since I follow Train (2009), we too begin with a focus on model specification.

Let me be clear: many students will see y and x later in this section and immediately think about “fitting” a model; that is, they assume they have data that they will put into an estimation routine to find estimates of the parameters of the model. We are not there yet! At this early stage in the book, we are only specifying models; that is, listing sets of assumptions about data generating processes and exploring the implications of those assumptions. There are no data yet. There will be parameters introduced in our choice of model specification, but we are not yet estimating those parameters. Have patience, we will eventually do these things.

1.1 Recap of Train Ch. 1

Train denotes the outcome in any given choice situation as y , determined by some observable factors collected in the vector \mathbf{x} and some unobservable factors collected in the vector $\boldsymbol{\varepsilon}$. The factors (\mathbf{x} and $\boldsymbol{\varepsilon}$) relate to the agent’s choice (y) through a function $y = h(\mathbf{x}, \boldsymbol{\varepsilon})$. We assume for the moment that we know $h(\cdot)$ and that \mathbf{x} and $\boldsymbol{\varepsilon}$ are length-one vectors (i.e., scalars) denoted x and ε .

Since we do not observe ε , we can’t predict y exactly. Instead, we focus on the probability of y , that is:

$$\begin{aligned} p(y|x) &= \Pr(\varepsilon \text{ such that } h(x, \varepsilon) = y) \\ &= \Pr(I[h(x, \varepsilon) = y] = 1) \\ &= \int I[h(x, \varepsilon) = y] f(\varepsilon) d\varepsilon \end{aligned} \tag{1.1}$$

For certain special choices of h and f , a closed-form expression¹ for the integral is available. But more generally, for almost any choice of h and f , we can approximate the integral through simulation. Train provides pseudo code on how to do so:

1. Repeat the following two steps many ($r = 1, \dots, R$) times:
 - Draw $\varepsilon^{(r)}$ from $f(\varepsilon)$.
 - Determine whether $h(x, \varepsilon^{(r)}) = y$. If so, set $I^{(r)} = 1$; else set $I^{(r)} = 0$.
2. Average the R values of I

Next we look at two examples where we use this procedure to approximate the $p(y|x)$ integral. The first example I made up. The second example is the binary logit model discussed by Train, but for which I'll show you how to simulate the $p(y|x)$ integral.

1.2 A Simple Example

Let's first set up a toy example to demonstrate how simulation can approximate the $p(y|x)$ integral. Suppose $x = 0.5$ and ε is uniformly distributed between -1 and 1. Define $h(x, \varepsilon)$ to be:

$$h(x, \varepsilon) = \begin{cases} 0 & \text{if } x + \varepsilon < 0 \\ 1 & \text{if } x + \varepsilon \in [0, 1] \\ 2 & \text{if } x + \varepsilon > 1 \end{cases} \quad (1.2)$$

We'll focus on the outcome $y = 2$. You can probably intuit that the $p(y = 2|x) = 0.25$ since only one quarter of the time will ε be sufficiently positive to make $x + \varepsilon > 1$.² Nevertheless, let's approximate the integral representation of $p(y = 2|x)$ through simulation to ensure we understand the process.

To walk you through the code, we first set a seed so that the pseudo-random numbers generated by `runif()` can be replicated exactly each time the code is run (even on different computers). We then specify that we will use 1,000 draws in the simulation and we create a vector `I` to hold our results. The simulation occurs via a `for()` loop where each time through the loop we take a draw of ε , calculate $0.5 + \varepsilon$ and check whether that sum is greater than one. If so, then $h(x, \varepsilon) = 2$ matching the value of y for the choice probability we want to assess — i.e., $p(y = 2|x)$ — and thus we store a 1 in the r^{th} position of `I`; otherwise we store a 0. We then average the values in `I` to get our approximation of $p(y = 2|x)$.

¹In this context, a *closed-form expression* means a way of writing the integral so that the anti-derivative sign is not part of solution. For example, the integral $\int x \, dx$ has the closed for expression $x^2/2$ plus some constant. We will see later that the Extreme Value distribution is often chosen for f predominantly because it leads to a closed for expression for the choice probability $p(y|x)$.

²More precisely, $p(y = 2|x) = \Pr(x + \varepsilon > 1|x = 0.5) = \Pr(\varepsilon > 0.5) = \int_{0.5}^1 f(\varepsilon) d\varepsilon = (0.5\varepsilon)|_{0.5}^1 = 0.25$.

```

set.seed(1234)

R <- 1000
I <- vector(length=R)

for(r in 1:R) {
  eps <- runif(1, min=-1, max=1)
  h <- 0.5 + eps
  I[r] <- as.integer(h > 1)
}
mean(I)

```

```
[1] 0.258
```

The simulated value 0.258 approximates the exact value 0.25 and can be made closer by increasing the number of draws used in the simulation.

Notice that we took draws of ε to empirically approximate the integral of $p(y|x)$, but we should not think of these draws as “data” in the sense of a dataset. They are ancillary numbers generated from the distribution of ε that we use to approximate the integral.

R users will recognize that we can shorten the code by taking advantage of R’s vectorized functions and its conversion of boolean values to 0/1 when used in mathematical operations. Here is a shorter implementation of the simulation; whether it’s “better” code is a matter of preference.

```

set.seed(1234)
R <- 1000
mean( runif(R, min=-1, max=1) + 0.5 > 1 )

```

```
[1] 0.258
```

That’s it. If you can generate pseudo-random draws from the density f and you know h , approximating a choice probability by simulation might only require a handful of lines of code.

1.3 A Binary Logit Model Example

As an example of a model with a complete closed-form solution, Train provides the binary logit model. In this example, we’ll let \mathbf{x} and $\boldsymbol{\varepsilon}$ be vectors of length two, and we’ll specify h and f as follows:

The “binary” part refers to the aspect of the model whereby the decision maker does one of two things; they either take an action ($y = 1$) or not ($y = 0$). To tie this model into a framework of behavior, we start with a utility function U . In Train’s specific example, utility is specified as

$$U(\mathbf{x}, \boldsymbol{\beta}, \varepsilon) = \mathbf{x}'\boldsymbol{\beta} + \varepsilon \quad (1.3)$$

where \mathbf{x} is a vector of observable explanatory variables, $\boldsymbol{\beta}$ is a vector of parameters that through the functional form $\mathbf{x}'\boldsymbol{\beta}$ effectively serve as weights on the covariates, and ε is a scalar index collecting the value of information used by the decision maker but unobserved to the researcher. Notice that we’re allowing \mathbf{x} and $\boldsymbol{\beta}$ to be vectors in this example, whereas they were scalars (or equivalently length-one vectors) in the previous example.

In this model, the threshold for action is 0 because the decision maker takes action ($y = 1$) when utility is positive; conversely, when utility is negative, the decision maker elects not to take the action (i.e., $y = 0$). Therefore we can specify h as:

$$h(\mathbf{x}, \boldsymbol{\beta}, \varepsilon) = \begin{cases} 0 & \text{if } U(\mathbf{x}, \boldsymbol{\beta}, \varepsilon) \leq 0 \\ 1 & \text{if } U(\mathbf{x}, \boldsymbol{\beta}, \varepsilon) > 0 \end{cases} \quad (1.4)$$

The “logit” part of the model’s name refers to the choice of f . The binary logit model assumes f is the logistic distribution:

$$f(\varepsilon) = \frac{e^{-\varepsilon}}{(1 + e^{-\varepsilon})^2} \quad (1.5)$$

Having specified h and f , let’s choose some values for \mathbf{x} and $\boldsymbol{\beta}$ and use simulation to approximate the integral for $p(y|\mathbf{x}, \boldsymbol{\beta})$.³ Let’s pick $\mathbf{x} = (0.5, 2)$ and $\boldsymbol{\beta} = (3, -1)$ such that $\mathbf{x}'\boldsymbol{\beta} = (0.5)(3) + (2)(-1) = 1.5 - 2 = -0.5$. We know from the closed-form solution to this model provided by Train that, with these values of \mathbf{x} and $\boldsymbol{\beta}$, the probability the decision maker takes action is:

$$p(y = 1|\mathbf{x}, \boldsymbol{\beta}) = \frac{e^{\mathbf{x}'\boldsymbol{\beta}}}{1 + e^{\mathbf{x}'\boldsymbol{\beta}}} = \frac{e^{-0.5}}{1 + e^{-0.5}} = 0.3775407 \quad (1.6)$$

We can approximate this integral as before. Below I use the function `rlogis()` to take $R=1000$ draws from the binary logistic distribution, and I approximate the integral with the proportion of times $\mathbf{x}'\boldsymbol{\beta} + \varepsilon$ is greater than the threshold for action (0):

³Only at the early stage of simulating data from a model to better understand the model do we pick values for \mathbf{x} and $\boldsymbol{\beta}$. Typically, \mathbf{x} will part of the data you collect and $\boldsymbol{\beta}$ will be parameters whose values you seek to estimate.

```
set.seed(2345)
R <- 1000

x <- c(0.5, 2)
beta <- c(3, -1)

U <- as.vector(x %*% beta) + rlogis(R)
mean(U > 0)
```

```
[1] 0.364
```

Our simulated value 0.364 approximates the exact value of the integral 0.378.

1.4 Key Learnings

The key learning from this chapter is that with discrete choice models our focus is on the *probability* of the choice outcome y – that is, $p(y|x)$. This choice outcome y results from the joint distribution f of unobserved factors ε and the behavioral model h that relates y to \mathbf{x} (and ε). The probability of the choice outcome $p(y|x)$ can be written in closed form for only very special choices of f and h , but for almost any choice of f and h we can simulate $p(y|x)$, as the two examples in this chapter demonstrate.

2 Properties of Discrete Choice Models

add

Part II

Behavioral Models

3 Logit

add

4 GEV

add

5 Probit

add

6 Mixed Logit

add

```
`<!-- quarto-file-metadata: eyJyZXNvdXJjZURpciI6ImNoYXB0ZXJzIn0= -->`{=html}
```

```
```${=html}
```

```
<!-- quarto-file-metadata: eyJyZXNvdXJjZURpciI6ImNoYXB0ZXJzIiwiaW9va0l0ZW1UeXB1IjoieY2hhcHRlc
```

## 7 Variation on a Theme

add

## 8 Numerical Maximization

add

# **Part III**

## **Estimation**

## 9 Drawing from Densities

add

## 10 Simulation-Assisted Estimation

add

# 11 Individual-Level Parameters

add



## 12 Bayesian Procedures

add

## 13 Endogeneity

add

## 14 EM Algorithms

add

# References

- Abramovich, Felix, and Ya'acov Ritov. 2023. *Bayesian Statistics*. CRC press.
- Gelman, Andrew, John B. Carlin, Hal S. Stern, David B. Dunson, Aki Vehtari, and Donald B. Rubin. 2013. *Bayesian Data Analysis*. CRC press.
- Goldberger, Arthur S. 1991. *A Course in Econometrics*. Harvard University Press.
- Kennedy, Peter. 2008. *A Guide to Econometrics*. John Wiley & Sons.
- McElreath, Richard. 2018. *Statistical Rethinking: A Bayesian Course with Examples in r and Stan*. Chapman; Hall/CRC.
- Train, Kenneth E. 2009. *Discrete Choice Methods with Simulation*. Cambridge University Press. <https://eml.berkeley.edu/books/choice2.html>.
- Wickham, Hadley. 2019. *Advanced r*. 2nd ed. Chapman; Hall/CRC.
- Wickham, Hadley, Mine Cetinka-Rundel, and Garrett Grolmund. 2023. *R for Data Science*. 2nd ed. O'Reilly Media.

# A Appendix

add